

1 Source-only model selection and limits of fixed spectral-sharing pipelines
2 in cross-subject motor-imagery EEG decoding

3 Short title: Source-only selection and fixed spectral-sharing pipelines in motor-imagery EEG

4 Ziyuan Zhu*

5 College of Artificial Intelligence Medicine, Chongqing Medical University, Chongqing, China

6 *Corresponding author: zzy2630816871@gmail.com

7 Postal address: Jinyun Campus, Chongqing Medical University, No. 61, Daxuecheng Middle

8 Road, Shapingba District, Chongqing 401331, China

9

10 **Abstract**

11 We audited source-only training-duration selection and evaluated fixed shared mu/beta input
12 pipelines for cross-subject motor-imagery EEG decoding. An exploratory four-class study used
13 nine BNCI2014_001 participants, excluding both target sessions from each corresponding model
14 fit and duration selection. Separate frozen binary pipelines were evaluated on PhysioNet (109
15 people), Cho2017 (52) and Lee2019 offline-training runs (54, two sessions), without target fitting
16 or adaptation. We re-audited saved predictions and source partitions. A BNCI signal analysis,
17 specified after decoder outcomes were known, described baseline-relative mu/beta power and
18 associations with saved binary LOSO balanced accuracy. Mean-rank duration selection increased
19 internal balanced accuracy from 33.72% to 42.67% compared with mean-loss selection; S3 and S8
20 contributed 87.57% of the aggregate gain. A matched-runtime fixed-duration control improved
21 four of nine people despite a positive mean. Shared-minus-broad external differences were +0.573
22 percentage points (pp) on PhysioNet (95% person-bootstrap interval $[-0.097, 1.242]$), -1.513 pp on
23 Cho2017 $[-2.249, -0.804]$ and $+0.204$ pp on Lee2019 $[-0.515, 0.969]$. Cho2017 and Lee2019
24 Holm-adjusted permutation p-values were 0.000100 and 0.604. These complete-pipeline contrasts
25 used unequal source-selected durations. The BNCI signal analysis retained all 5,184 trials; its
26 2,592 hand-imagery trials gave mean signed laterality of -0.250 dB in mu and -0.168 dB in beta,
27 with signs varying between people. Training-duration choices yielded large but concentrated gains
28 in this development benchmark. Shared-input pipelines showed no consistent external advantage.
29 BNCI sensor patterns provide descriptive context; nine-person associations do not establish a
30 decoder mechanism. Unknown voltage calibration and offline evaluation limit deployment
31 conclusions.

32 **Introduction**

33 Motor-imagery EEG decoding must distinguish imagined actions in scalp recordings whose spatial
34 and spectral structure varies between people. A calibration-free model should transfer to a new

35 person without fitting to that person’s labels or distributional statistics. This requirement differs
36 from within-person evaluation and adaptation using unlabeled target recordings. BCI Competition
37 IV dataset 2a is a small, well-documented benchmark [1,2], but its nine participants make the
38 cohort mean especially sensitive to a few large changes. Reviews and multi-dataset evaluation
39 frameworks document variation between datasets and show how evaluation design shapes reported
40 performance [3,4].

41 A decoder that predicts predominantly one class can have balanced accuracy near chance even
42 when the same architecture performs well after a different training duration. This failure may be
43 attributed to an input distribution that requires normalization or a different representation,
44 overlooking the role of model selection. Likewise, an increase in average balanced accuracy may
45 reflect recovery in one or two participants rather than consistent improvement across people.

46 Source-only fitting prevents target leakage within each fit, but it does not provide independent
47 confirmation when the research programme was developed after examining the same benchmark.

48 We address three questions. First, can heterogeneous source-validation loss curves lead to very
49 short training schedules, and does within-fold rank aggregation change those schedules? Second,
50 how does a pipeline that processes predefined mu/beta views with one shared EEGNet compare
51 with broad input, alongside spatial, capacity, and robustness controls? Third, how do the complete
52 source-selected pipelines perform in external binary transfer? A separate descriptive analysis
53 characterizes measured BNCI sensor signals without changing the decoders. The external evidence
54 comprises an earlier PhysioNet evaluation [5,6] and a subsequently frozen, separately fitted
55 pipeline for Cho2017 [7,8] and Lee2019 [9,10]. The evaluations use different channel sets and
56 time windows and are reported separately.

57 Here, we audit source-only model selection as part of cross-subject decoder evaluation [11]. The
58 development analyses trace how training-duration decisions affect apparent model failure and
59 distinguish concentrated aggregate gains from participant-level benefit. Separate frozen external
60 evaluations retain adverse and uncertain outcomes, with participants as the unit of inference. A

61 downstream BNCI sensor analysis adds measured-signal context while keeping physiological
62 description separate from decoder attribution.

63 **Related Work**

64 Classical motor-imagery decoders use spatial and spectral structure. CSP estimates discriminative
65 spatial projections from class covariance structure [12], and filter-bank CSP combines spatial
66 features across frequency bands [13,14]. These methods inform the spatial and spectral controls
67 used here. The implemented filter bank is FBCSP-inspired rather than an exact reproduction; its
68 projections, feature scaling and classifiers are fitted on source partitions under the same target-
69 information restriction as the neural models.

70 EEGNet combines temporal filtering, depthwise spatial filtering and separable convolutions in a
71 compact architecture [15]; broader CNN work investigates learned EEG decoding and
72 visualization [16]. Our shared-input model processes predefined mu and beta views with one
73 configured EEGNet and averages their class logits. Its parameter count matches one branch, but
74 computation on two views and shared BatchNorm exposure differ. These bands define
75 experimental input views; they do not provide verified physiological source separation.

76 Algorithm reviews and multi-dataset benchmarks document performance variation between
77 methods, people and datasets [3,4]. Domain-generalization evaluation must also account for model
78 selection [11]. Here, fitting is restricted to source participants within a repeatedly used
79 development benchmark, followed by external evaluation of frozen pipelines. Methods that fit on
80 target labels or unlabeled target distributions have a different information budget and were not
81 compared directly.

82 Recent domain-generalization studies seek invariance across source domains for unseen-subject
83 motor-imagery decoding. EEG-DG aligns marginal and conditional distributions across multiple
84 source domains [17]. Zheng et al. combine spectral/task knowledge distillation with cross-source
85 correlation alignment and distance regularization [18]. These approaches pursue source-only
86 generalization without target adaptation, although a matched comparison still requires compatible

87 information budgets and selection procedures. Our contribution is an audit of source-only model
88 selection, participant heterogeneity and complete frozen-pipeline transfer, with adverse and null
89 results retained in a traceable evidence record. We do not introduce a new domain-generalization
90 algorithm or rank the evaluated pipelines against these methods.

91 EEG pretraining offers another route to transferable representations. LaBraM learns generic
92 representations from large EEG collections [19], while EEGPT uses pretrained transformers for
93 transferable EEG representations [20]. Fair comparison requires the pretraining data, downstream
94 fitting procedures and computational budgets to be specified. Neither model was evaluated here,
95 and equivalence to our target-information contract was not established; the findings concern the
96 selected source-only pipelines.

97 Event-related desynchronization and synchronization describe decreases and increases in spectral
98 power relative to a reference interval [21]. Earlier motor-imagery work investigated sensorimotor
99 rhythm changes [22]. These concepts guide the separate characterization of measured signals, for
100 which the baseline and estimator are specified explicitly. Scalp power patterns do not establish
101 anatomical source localization or demonstrate that a learned decoder used those patterns.

102 **Materials and Methods**

103 **Experimental Design**

104 The study comprises exploratory development and two external-evaluation stages. In the four-class
105 BNCI2014_001 development stage, both sessions of each target person are excluded from the
106 corresponding fitting and duration selection. Earlier benchmark outcomes informed later
107 candidates, so the nine-person comparisons remain exploratory. The earlier PhysioNet external
108 evaluation (Q14) applies separately fitted binary BNCI models to PhysioNet; a metadata-driven
109 sampling-rate amendment followed partial target execution. The frozen Cho2017/Lee2019 transfer
110 experiment (Q15) uses a new binary source pipeline with a different montage, crop and reference,
111 evaluating its frozen broad, shared-input and CSP models on Cho2017 and Lee2019 offline-

112 training segments. These datasets are external to fitting, but participant identity nonoverlap across
113 providers was not independently established.

114 The primary external estimand is the mean paired balanced-accuracy difference between the
115 frozen shared-input and broad pipelines, with equal weight for each person and separate estimates
116 for each cohort. The PhysioNet evaluation uses 17 shared and 18 broad epochs; the frozen
117 Cho2017/Lee2019 evaluation uses 19 shared and 14 broad epochs, all selected on source data.
118 Each contrast therefore evaluates complete selected pipelines rather than the isolated effect of
119 sharing. CSP provides context, and neither external stage compares rank with mean-loss selection.
120 The descriptive BNCI physiological analysis (Q16) was defined after decoder results were known
121 and uses BNCI signals only. Its raw-signal summaries and associations provide no feedback into
122 the decoders and cannot establish their physiological mechanism.

123 The corresponding saved summary is presented in Fig 1.

124 Fig 1. Study workflow and source–target information boundaries. Study workflow and information
125 boundaries. Both sessions of each internal target person are excluded from the corresponding
126 source fit; these development analyses remain exploratory. Q14 and Q15 use separate binary
127 models and preprocessing contracts. At inference, deterministic adapters and frozen source-
128 derived transformations are applied without fitting, adaptation or BatchNorm updates. Event and
129 class metadata establish eligibility and canonical label mapping; ground-truth labels are used for
130 scoring. The separate BNCI-only physiological analysis reads native signals and labels after
131 decoder outcomes are known, without feeding back into fitting or selection. Results are aggregated
132 by person. Q14's post-partial-run amendment and the external differences in selected duration are
133 described in the Methods and comparison-fairness table.

134 **Datasets, Populations and Task Separation**

135 BNCI2014_001 corresponds to BCI Competition IV dataset 2a [1,2]. Nine participants completed
136 two recording sessions, each containing six runs of 48 trials: 12 each of left-hand, right-hand, feet,
137 and tongue imagery. Twenty-two EEG channels and three EOG channels were sampled at 250 Hz.

138 Acquisition used a left-mastoid reference, right-mastoid ground, 0.5–100 Hz band-pass, and 50-Hz
139 notch; the offline operations below were applied in addition to this acquisition processing. The
140 primary neural input contains EEG channels only. The archived metadata contain 5,184 trials, 576
141 per participant and 1,296 per class. All 488 expert artifact-flagged trials remain in the primary
142 evaluation population. Including all trials differs from the original competition’s artifact-free
143 scoring; source-clean conditions exclude flagged source-training trials but retain every target trial.
144 The external EEG Motor Movement/Imagery Database version 1.0.0 [5] provides a separate binary
145 task. We include all 109 subjects and only imagery runs 4, 8, and 12; T1 and T2 denote left- and
146 right-fist imagery. Executed-movement, bilateral-task, and rest events are excluded from this
147 contrast. The final audited inventory contains 327 EDF files verified against official checksums
148 and 4,918 unique trials. Although the acquisition description specifies 160-Hz sampling, the file
149 headers identify three subjects with 128-Hz recordings. Four-class and binary balanced accuracy
150 have different chance levels and are reported separately.

151 Cho2017 [7,8] provides left/right-hand imagery recordings from 52 people at a native sampling
152 rate of 512 Hz. The audited evaluation retains 10,520 imagery trials: 49 people have 200 trials and
153 three (S7, S9, S46) have 240. The Lee2019 motor-imagery data [9,10] contain 54 people recorded
154 at 1,000 Hz in two sessions. We retain only EEG_MI_train, the provider’s offline-training runs,
155 giving 100 trials per session and 10,800 trials in total. EEG_MI_test, other paradigms, and non-
156 imagery events are excluded. The provider’s term “train” identifies a recording segment; these
157 external trials were used only for inference and scoring, never for parameter fitting. The three
158 external cohorts cannot be assumed to represent 215 independently verified unique identities
159 across providers.

160 The corresponding saved summary is presented in Table 1.

161 Table 1. Distinct populations and computational endpoints.

Analysis	People	Classes	Unique trials	Primary tensor
Internal BNCI	9	4	5,184	22×750 at 250 Hz
Q14 binary BNCI source	9	2	2,592	22×480 at 160 Hz
Q14 external PhysioNet	109	2	4,918	22×480 at 160 Hz
Q15 binary BNCI source	9	2	2,592	21×320 at 160 Hz
Q15 external Cho2017	52	2	10,520	21×320 at 160 Hz
Q15 external Lee2019	54	2	10,800	21×320 at 160 Hz

162 Both sessions of each BNCI outer target are held out. Q14 and Q15 have separate binary source
163 fits and preprocessing contracts, and neural inputs contain EEG channels only. Prediction rows
164 repeat trials across models and seeds; they are not independent participants. Lee sessions are
165 pooled within person for the primary endpoint.

166 **Strict Zero-calibration Information Contract**

167 Within each evaluation fit, strict zero-calibration means that no parameter, state, preprocessing
168 coefficient, hyperparameter, duration, seed, decision threshold or model choice is fitted or selected
169 using the current target person’s recordings or labels. Neural weights and BatchNorm statistics,
170 shallow-model coefficients, source-fitted transforms and source-selected schedules are fixed
171 before target inference. Target data are not used to fit normalization, covariance alignment or
172 adaptation, update BatchNorm, stop training or select seeds. This restriction applies to each fit; it
173 does not make the repeatedly used BNCI development programme independent of earlier target
174 outcomes.

175 Target EEG provides model inputs; event and class metadata determine declared eligibility and
 176 canonical label mappings. Labels are then used to score saved predictions. Cho recordings are
 177 stored in class-specific containers, so the adapter has access to native class metadata. Fixed
 178 filtering, channel selection, resampling and common-average referencing operate on target signals
 179 without fitting a target-dependent decoder. The Q15 Helmert basis depends only on channel count;
 180 scalars, CSP/PCA, whitening and EOG-regression coefficients, where used, are estimated on
 181 source-training data and applied unchanged. Separate physiological summaries use raw signals and
 182 labels but provide no feedback into bands, crops, weights, thresholds, schedules or model choice.
 183 The absence of target fitting does not establish physical acquisition calibration. Q15's native-
 184 numeric-as-microvolt convention does not verify the exported voltage units. Earlier within-person
 185 and cross-session P2 models use target-person labelled calibration and remain outside the strict
 186 zero-calibration endpoints. The source-fitted P4 EOG regression requires synchronous target EOG
 187 and thus additional sensors, although it fits no target coefficients.
 188 The corresponding saved summary is presented in Table 2.
 189 Table 2. Target information and fitting boundaries.

Operation	Information read	Target fitting	Interpretation
Eligibility, fixed channel and class maps, cue indexing	Native rate/channel/event/class metadata	None	Declared metadata adapter
Filtering, resampling, crop and Q15 CAR	Target EEG, fixed native metadata	None	Deterministic offline operations
Q15 Helmert coordinates	Channel count and signal	None	Analytical CAR-subspace basis
Source-fitted scaler/CSP/PCA/whitenin g/EOG regression	Source data to fit; target signal to apply	None	Frozen source coefficients; EOG needs extra sensors

Operation	Information read	Target fitting	Interpretation
Neural inference	Target tensors	None; weights and BatchNorm fixed	Strict zero-calibration prediction
Scoring and separate Q16 characterization	Predictions/raw signals and canonical labels	None	Downstream analyses with no decoder feedback
P2 within-person or cross-session models	Target-person labeled training trials	Yes	Calibrated context; separate endpoint

190 The restriction applies to fitting and selecting the corresponding decoder. It permits deterministic
191 computation on target signals, metadata inspection, and label use for scoring and separate
192 physiology. These boundaries do not make the development history outcome-blind.

193 **Offline Preprocessing and Source-only Partitions**

194 For internal BNCI analyses, fourth-order zero-phase Butterworth filters are applied within each
195 native recording run before epoch extraction. The half-open interval [2.5,5.5) s after the MAT trial
196 start corresponds to [0.5,3.5) s after the cue and yields 750 samples. The primary neural pipeline
197 applies no baseline subtraction, additional notch, or additional re-reference. MNE-loaded values in
198 volts are multiplied by 10^6 to give network inputs in microvolts. These offline zero-phase
199 operations use future samples, so this pipeline does not constitute a causal online decoder.

200 The broad mean-loss and mean-rank EEGNet conditions use 4–40 Hz input. Restricted broad input
201 uses 8–30 Hz, and the shared representation uses fixed mu (8–13 Hz) and beta (13–30 Hz) views.
202 These predefined bands are experimental choices; they are neither participant-specific estimates
203 nor physiological source separation. Any fitted scaler, CSP/PCA projection, or whitening
204 transform is estimated only on the applicable source-training partition.

205 Each primary four-class outer split holds out both sessions of one participant. The eight source
206 people contribute 4,608 trials and the target contributes 576. Sorted source IDs define four
207 consecutive, two-person inner-validation groups. Each inner fit trains on six people/3,456 trials
208 and validates on two people/1,152 trials. Thus, source training and validation are separated by

209 person. Final evaluation uses three fixed seeds, with scores averaged within person before
210 aggregation across people. BatchNorm statistics remain fixed during model evaluation and are not
211 updated by target examples.

212 **Training-duration Selection and Diagnostic Controls**

213 The mean-loss rule selects the earliest epoch that minimizes validation cross-entropy averaged
214 with equal weight across the four inner folds. Candidate epochs are 1–40. The mean-rank rule
215 ranks epochs within each fold by validation cross-entropy, assigns average ranks to ties, averages
216 the ranks across folds, and selects the earliest minimum. This rule is invariant to strictly increasing
217 (order-preserving) transformations of the losses within each fold, although that property does not
218 guarantee better target prediction. The rank follow-up reuses the 36 frozen original inner
219 trajectories and trains 27 new final models after the schedule has been frozen.

220 Selected epoch = earliest argmin over e of $(1/K) \sum_k \text{rank}_k[L_k(e)]$.

221 The rank candidate was motivated by earlier results on the same BNCI participants, so the internal
222 comparison remains exploratory despite source-only selection within each fit. A four-cell
223 diagnostic for S3 crosses raw and source-normalized input with two and 16 fixed epochs. Here,
224 raw input means filtered, microvolt-scaled EEG without fitted source normalization. Two cells
225 reuse the archived original conditions and two are new diagnostic conditions preserved in the
226 archive; the comparison involves one participant and does not provide population-level replication.

227 Source normalization estimates per-channel mean and scale from source-training samples only.

228 Later controls compare fixed 20-epoch training with raw-loss-selected schedules. A matched-

229 runtime experiment retains the historical source-selected epoch counts but retraining the 27 final

230 models using the same runtime and raw files as the fixed-duration comparator. Epochs are not

231 reselected under the newer runtime. Source-count controls use $k = 2, 4$, or 6 source participants,

232 four predefined identity subsets per count, and the $k = 8$ anchor; all counts use fixed 20 epochs.

233 The 0train-versus-1test source-session contrast has equal training budgets: 2,304 trials, 20 epochs,

234 and 720 optimizer updates per fit. Source-count contrasts and comparisons of one source session

235 with the both-session anchor change both the number of training examples and optimizer updates,
236 preventing attribution to a single causal factor.

237 The population raw-versus-source-normalized Q5/Q6 comparison also changes the selected
238 duration for S3 from 2 to 16 epochs and S5 from 13 to 12. Its aggregate difference therefore
239 cannot be attributed solely to normalization. The fixed-duration Q7 four-cell diagnostic separates
240 normalization and duration more directly, but only for S3.

241 **Neural Representation and Baseline Fairness**

242 The four-class model is Braindecode 1.5.1 EEGNet with 22 input channels, 750 samples, $F1 = 8$,
243 depth multiplier $D = 2$, $F2 = 16$, temporal kernel length 64 samples, and dropout 0.25. It has 2,932
244 trainable parameters in the archived implementation. Training uses Adam, learning rate 0.001,
245 batch size 64, cross-entropy, zero weight decay, and no scheduler or default augmentation. The
246 selection seed is 20260923 and the three final seeds are 20260924, 20260925, and 20260926. This
247 is a configured EEGNet study, not an exact reproduction of the original paper's preprocessing or
248 splits [15].

249 For shared mu/beta input, both band views pass through the same EEGNet in one concatenated
250 mini-batch. Their class logits are averaged before cross-entropy during training and before softmax
251 during inference. Network parameters and BatchNorm are shared, while dropout acts per view.

252 The model has no learned band weights, attention mechanism, or independently trained branches.
253 Its parameter count matches one EEGNet, but forward-pass computation and preprocessing differ.
254 Shared prediction: $p = \text{softmax}[(f\theta(X\mu) + f\theta(X\beta))/2]$.

255 Controls include individual frequency bands; fixed source-selected durations; Welch log-
256 bandpower with source-fitted shallow classifiers; eight source-only CSP or PCA time-series
257 projections per band; log-Euclidean covariance features; source-only normalization; and source-
258 clean training. The CSP/PCA neural controls both use 16 projected channels and source-fitted
259 standardization. Their input widths therefore match each other but differ from the 22-channel

shared-logit model. Architectural ablations use two independent networks, early band stacking, four shared bands, and a broadband capacity comparator. The independent two-band model has 5,864 parameters and its comparator has 5,914, a 0.85% mismatch. The traditional filter-bank implementation is FBCSP-inspired rather than a literal reproduction of Ang et al. [13,14].

The corresponding saved summary is presented in Table 3.

Table 3. Matching and remaining differences in decoder comparisons.

Comparison	Matched aspects	Remaining difference / scope
Q5 mean-loss vs Q8 mean-rank	Nine-person all-trial LOSO, 22×750, 4–40 Hz, original 36 inner trajectories, final seeds	Schedule changes; candidate arose after BNCI outcomes; exploratory
Q5 raw vs Q6 source normalization	Same all-trial targets and EEGNet family; source-only scaler fit	Selected epochs also change: S3 2 → 16, S5 13 → 12; not normalization alone
Q7 S3 raw/normalized ×2/16 epochs	Epochs fixed within each raw/normalized pair	One-person diagnostic
Q9 broad vs shared, reselected durations	Same targets, dimensions, EEGNet parameter count, optimizer/seeds	Bands, two-view computation, BatchNorm exposure and schedules differ
Q9-E002 fixed-schedule broad/shared	Same targets and frozen Q8 per-target epochs	Duration-controlled sensitivity; views/computation still differ
Q10 CSP vs PCA neural projections	Eight source-fitted components per band; 16 projected channels and source scaling	Matched to each other, not to 22-channel shared-logit input
Q11 independent branches vs capacity comparator	Same targets; approximate trainable capacity	5,864 vs 5,914 parameters; branch computation/normalization differ
Q12 GroupDRO vs balanced ERM	Same runtime and batching/sampling policy	Declared objective comparison

Comparison	Matched aspects	Remaining difference / scope
Q12 whitening/augmentation vs historical Q8	Same nominal target population and source-only rule	Historical baseline uses another runtime
Q13 fixed20 vs matched-runtime CE schedule	Same actual runtime, raw files and three seeds	Historical epochs replayed, not reselected in the newer runtime
Q13 source-session 0train vs 1test	2,304 source trials, 20 epochs, 720 updates	Session composition changes; both-session anchor also changes exposure
Q14 external broad vs shared	109 people/4,918 trials; 22×480; same source seeds and target rule	18 vs 17 epochs; metadata amendment after partial execution
Q15 external broad vs shared	Identical people/trials within cohort; 21×320 adapter, neural family/seeds	14 vs 19 epochs and one-view vs two-view processing
Q14/Q15 CSP4+LDA context	Same cohort scoring population; source-only fit	One shallow model vs three neural seeds; Q15 Helmert; differs from Q4 CSP8/OAS

Each comparison controls the aspects listed, while retaining the stated implementation differences. Q14/Q15 therefore compare complete frozen pipelines. Matching selected aspects does not isolate every causal factor, and the CSP comparisons remain descriptive.

Source-only Robustness Interventions

The six completed robustness conditions are source-pooled whitening, source-balanced empirical risk minimization (ERM), source-group distributionally robust optimization (GroupDRO), channel dropout, gain perturbation, and their combination. Source whitening uses equal-person covariance aggregation, a ridge of 5% of the mean eigenvalue, and scale preservation; no target covariance or alignment is fitted. GroupDRO treats each source person as a group [23]. Persistent log weights increase by 0.05 times detached per-group cross-entropy, are normalized by log-sum-exp, and weight the group loss. Its matched comparator is balanced ERM, which preserves the batching/sampling intervention. Training-only channel dropout uses probability 0.1 with

278 surviving-channel scale 1/0.9; gain perturbation uses $U[0.8,1.2]$ and precedes dropout in the
279 combined condition. Validation and target inputs are not augmented.
280 Whitening and augmentation conditions are compared with the previously frozen broad rank-
281 selected baseline. The historical baseline ran on a different runtime from the robustness batch, so
282 these differences reflect both intervention and runtime variation. GroupDRO versus balanced
283 ERM provides a within-runtime comparison of training objectives. This distinction accompanies
284 the numerical table and limits attribution of cross-batch differences to the named interventions.

285 **Earlier Binary Source Freeze and PhysioNet Evaluation (Q14)**

286 The binary BNCI source population contains 2,592 left/right trials, 288 per person. Filtering and
287 epoching occur at 250 Hz before deterministic polyphase resampling by 16/25 to 160 Hz, giving
288 22×480 samples. The broad binary model uses 8–30 Hz; the shared model uses the fixed mu/beta
289 views. Source-only duration selection across all nine participants uses validation groups
290 comprising people 1–2, 3–4, 5–6, and 7–9, with equal fold weighting. The frozen training
291 durations are 18 epochs for the broad model and 17 for the shared model. Checkpoints, channels,
292 labels, contrast, seeds, and preprocessing are fixed before external access. The source-only CSP4 +
293 LDA comparator uses $\text{reg} = \text{None}$ and LDA solver = svd; it differs from the internal CSP8
294 OAS/shrinkage-LDA configuration. The external neural comparison evaluates complete pipelines
295 with different source-selected durations and therefore cannot isolate parameter sharing.
296 A partial external run stopped at S88 because its 128-Hz header violated the original 160-Hz
297 assumption. A versioned amendment used metadata from the remaining headers without
298 consulting predictions or scores. S88, S92, and S100 are filtered and epoched at their native rate,
299 then resampled by 5/4 from 384 to 480 samples. The other 106 subjects retain native 160-Hz
300 processing. The 87 completed subject outputs from the earlier run are copied byte-for-byte, and 22
301 additional subjects complete the cohort. The amendment introduces no target statistics, fitted
302 normalization, adaptation, seed selection, or class-performance-dependent processing. Because it

303 followed partial external execution, the combined result does not represent an unmodified
304 execution of the original freeze.

305 The original aggregate validation failed because it applied an inappropriate exact floating-point
306 equality check to re-serialized CSV probabilities. A separate cross-host runtime mismatch later
307 prompted a versioned portability amendment. Both limitations are documented. The versioned
308 validation-only audit checked official EDF checksums and event identities and recomputed saved
309 probability/argmax consistency, subject metrics, and statistics. It passed with zero new fits and
310 zero new model-inference rows. It therefore validates saved outputs; it does not regenerate all
311 external predictions from checkpoints on the later validation host.

312 **Frozen Cho2017/Lee2019 Transfer (Q15)**

313 The frozen Cho2017/Lee2019 study is a separate binary transfer experiment using different
314 checkpoints from Q14. All nine BNCI participants supply 2,592 left/right trials. The common
315 ordered montage comprises Fz, FC3, FC1, FC2, FC4, C5, C3, C1, Cz, C2, C4, C6, CP3, CP1, CPz,
316 CP2, CP4, P1, Pz, P2, and POz. Each trial is processed in a fixed native-rate context from -1.5 to
317 $+4.5$ s relative to the cue. Common-average reference is computed across these 21 selected
318 channels, followed by native-rate zero-phase second-order-section Butterworth filtering with
319 prototype order four. Broad input uses 8–30 Hz; the shared views use 8–13 and 13–30 Hz.

320 Deterministic polyphase resampling yields 160 Hz, followed by the half-open cue crop $[0.5, 2.5)$ s,
321 equivalent to samples $[320, 640)$ of the resampled six-second context. The resulting model tensor is
322 21×320 . These preprocessing and input differences preclude interpreting a comparison with Q14
323 as an effect of external provider alone.

324 The Q15 adapter interprets native numeric values as microvolts by convention. Physical voltage
325 calibration, the original Cho hardware reference, and hardware cue latency have not been
326 independently verified. The first 64 of the Cho payload's 68 rows are EEG; four EMG rows are
327 excluded. Channel mapping follows the publication's numbered montage because the payload
328 does not supply names. Cho native labels left = 1 and right = 2 are retained. Lee native labels left =

2 and right = 1 are deterministically mapped to canonical left = 1 and right = 2. Lee MATLAB cue indices are converted using $t - 1$; the provider's stored segments align with Python t , leaving a disclosed one-sample (1-ms) indexing discrepancy that was not changed using outcome scores. The supplement documents these adapter assumptions and the file-level checks.

Source-only duration selection uses four groups comprising people 1–2, 3–4, 5–6, and 7–9, candidate epochs 1–40, equally weighted mean within-fold cross-entropy ranks, and the earliest minimum. Final durations are 14 broad epochs and 19 shared epochs. The selection seed is 20260923, and final seeds are 20260924–20260926. Training uses the same configured EEGNet family, optimizer, and fixed-band logit averaging described above, with the Q15 channel/time dimensions. A source-fitted CSP4 + LDA comparator applies a fixed 21-to-20 Helmert coordinate transform after common-average reference, followed by source-only fitting. The external neural comparison evaluates two complete frozen pipelines with different source-selected durations and cannot isolate parameter sharing or band decomposition alone. Source training comprises eight inner neural fits, six final neural fits, and one shallow fit: 15 in total. Migration reused these artifacts without additional source or target fitting.

Official raw metadata for 18 BNCI files and all declared Cho/Lee files were audited before the preprocessing specification was committed; its pre-fit publication was recorded at 2026-10-04 01:33:03 UTC. Source training was completed at 01:52:42 UTC. The separate inference freeze was committed on 2026-10-05 at 09:31:05 UTC, before external prediction, fixing channels, timing, label adapters, transformations, checkpoints, seeds, scoring units, and primary statistical procedures. Event metadata may be inspected to determine preprocessing eligibility, but external labels are not used for fitting or model selection. BatchNorm remains in evaluation mode; there is no target alignment, normalization fit, adaptation, or post hoc seed selection. These repository-defined freezes followed external metadata inspection and do not constitute independently registered preregistration.

354 All 52 Cho2017 and 54 Lee2019 participants were processed under the declared adapters. The
355 archived independent computational validator replays processing from raw data to epochs and
356 predictions from the frozen models, checks canonical events and scoring populations, verifies links
357 to inference and source artifacts, and recalculates aggregate statistics. Validation passed subject to
358 the stated acquisition-calibration limitations. The supplementary reproducibility documentation
359 identifies the exact input, parameter, checkpoint and result revisions.

360 **BNCI-only Physiological Signal Characterization (Q16)**

361 The BNCI characterization separately describes native BNCI2014_001 recordings: 18 files, 108
362 motor-imagery runs, nine people and both sessions, retaining all 5,184 four-class trials and their
363 artifact flags in the primary population. The signals retain all 22 EEG sensors at native 250 Hz,
364 with no additional bandpass, notch, resampling, reference transform, source scaler or learned
365 projection. Three EOG channels are excluded. EEG column order follows the provider montage;
366 individual MAT files do not contain channel-name strings. Native one-based trial starts are
367 converted to zero-based samples; the cue is 500 samples after the trial start. The fixed cue-relative
368 baseline is $[-1.5, -0.5]$ s and the task window is $[0.5, 2.5]$ s. Before power calculation, all originals
369 are rehashed against archived source provenance, and labels, artifact flags, finite samples, montage
370 order and complete window bounds are checked. All 18 source SHA-256 digests, 5,184 event
371 identities and 228,096 trial-channel-band eligibility rows passed these checks; no trial was
372 excluded by a power guard. The primary population retained all 488 flagged trials, including 246
373 of the 2,592 hand trials.

374 Welch power is defined separately within each trial's baseline and task windows using float64
375 native numeric values, a periodic Hann taper, 250-sample segments, 125-sample overlap, 250-
376 point FFT, linear detrending, one-sided density and arithmetic segment averaging. This gives 1-Hz
377 bins. Closed 8–13-Hz and 13–30-Hz masks are integrated with the trapezoidal rule in physical
378 frequency; the 13-Hz boundary supplies half-bin weights to the adjacent integrals. The trial-level
379 change is $10 \log_{10}(P_{\text{task}}/P_{\text{baseline}})$, in dB. Nonfinite or nonpositive powers require an explicit

380 invalid-row receipt rather than an epsilon or clipped ratio. The 1-s baseline contributes one Welch
 381 segment and the 2-s task contributes three overlapping segments. Unequal log-power estimator
 382 variance can create a positive offset even without a true power change, so these dB summaries are
 383 not claimed to be unbiased ERD estimates.

384 Arithmetic mean trial dB is aggregated within person, session, class, channel and band, followed
 385 by equal averaging of the two sessions. This averages log ratios rather than taking the logarithm of
 386 averaged powers. Left/right-hand C3/C4 summaries use common trial eligibility. The signed hand
 387 descriptor is $L = [(C3_{\text{right}} - C4_{\text{right}}) + (C4_{\text{left}} - C3_{\text{left}})]/2$. A negative descriptor means the
 388 contralateral task/baseline change is lower than the ipsilateral change; it does not alone establish
 389 absolute contralateral ERD. C3, Cz and C4 absolute changes, separate hand terms, full scalp
 390 summaries, sessions and artifact strata are retained to show that distinction.

391 The physiological parameters were fixed on 6 October 2026 after the decoder outcomes were
 392 known and before the physiological calculations. The analysis is therefore descriptive and does not
 393 provide independent confirmation. Associations use existing matched-binary Q14-E001 BNCI
 394 LOSO scores, averaged over the three fixed neural seeds within each of the same nine people; all-
 395 nine-source fits and external-cohort scores are excluded. The broad EEGNet score is the
 396 designated descriptive endpoint. Six Spearman coefficients cover the three fixed models and two
 397 fixed bands, with no p-values, tertile groups or outcome-selected bands/models. Decoder imagery
 398 covers [0.5,3.5) s, whereas the physiological power analysis covers [0.5,2.5) s. No external
 399 physiology was calculated, and no physiological output was used to fit, select or update a decoder.

400 The parameter freeze is commit 050e01b028aaab8e3d745934b13b2d17e9bb0a7a; its public
 401 readback preceded raw-power calculation. The preprocessing freeze, raw-data receipt, execution
 402 manifest and final independent validation link the specification and inputs to the completed tables
 403 in Q16-P001-BNCI-20261006.

404 **Statistical Analysis**

405 Balanced accuracy = $(1/C) \sum_c TP_c / (TP_c + FN_c)$.

406 The held-out person is the unit of statistical inference. Paired seeds are averaged within person;
407 predefined source-identity subsets are averaged within person and source count. Shallow
408 conditions have one prediction per trial rather than three independent neural seeds. Paired
409 differences are expressed in balanced-accuracy percentage points (pp). Person-bootstrap percentile
410 intervals resample people, not trials, sessions, seeds, or source subsets. External intervals describe
411 variation across target people conditional on the frozen source models; they do not quantify
412 source-cohort sampling uncertainty. Internal analyses remain exploratory because benchmark
413 outcomes informed the research program and outer training sets overlap. For Lee2019, the two
414 retained sessions are pooled before calculating each person/seed balanced accuracy. The three
415 neural-seed scores are then averaged within person, giving an average of scores rather than an
416 ensemble of probabilities. A single CSP score is retained per person. Cross-cohort score
417 differences are descriptive because Q14 and Q15 differ in source preprocessing and checkpoints.
418 For each archived experiment, we report its original statistics and the interval resampling
419 procedure. Rank-selection intervals use 20,000 person-bootstrap draws with seed 20260923. The
420 archived shared-minus-broad representation contrast and robustness intervals use 20,000 draws
421 with seed 20260924. Source-composition and matched-runtime intervals use 10,000 draws with
422 seed 20260926 and exhaustive two-sided sign flips over nine participant differences; familywise
423 exploratory p-values use Holm correction within the frozen families. The earlier PhysioNet
424 primary contrast uses 20,000 draws with seed 20260924 and an exact two-sided binomial sign test
425 excluding exact zero differences. Its primary hypothesis is shared-minus-broad; contextual
426 comparisons with CSP do not replace it. A supplementary numerical audit uses 200,000 draws
427 with seed 20261002 and is labelled separately. An interval excluding zero need not agree with
428 every small-sample test. For Q15, the predeclared family contains exactly two shared-minus-broad
429 contrasts, one for each external cohort. Percentile person-bootstrap intervals use 20,000 draws with
430 seed 20260924 [24]. Two-sided person-level sign-flip tests use 20,000 Monte Carlo draws with
431 seed 20261003, assuming sign exchangeability of paired person effects under the null, evaluating

432 the absolute mean, and applying the $(\text{exceedances} + 1)/(20,000 + 1)$ correction [25]. Holm
433 adjustment controls this two-contrast family [26]; it does not retroactively include the earlier
434 PhysioNet or internal hypotheses. The smallest attainable Monte Carlo p-value is $1/20,001$. The
435 intervals are unadjusted 95% intervals, and CSP comparisons are descriptive context. Descriptive
436 Q15 counts of positive, negative, and tied person effects use an absolute 10^{-15} tolerance for ties;
437 the frozen primary bootstrap and sign-flip calculations retain the saved numerical effects.
438 One historical rank-selection script labelled a nine-person binomial calculation as a sign test while
439 retaining two unchanged participants in its denominator. We report its $p = 0.1797$ as an archived
440 nine-person sign sensitivity. A conventional tie-excluding calculation uses seven nonzero changes
441 and gives $p = 0.015625$; the paired t-test gives $p = 0.111$. These procedures are reported separately
442 rather than treated as interchangeable or used to select a favourable result. Multiplicity corrections
443 from later declared families are not retrospectively applied to unrelated earlier statistics.
444 The physiological characterization was specified after decoder outcomes were known and remains
445 descriptive. Physiological aggregation and decoder associations use nine person-level records;
446 channels, trials, sessions and seeds are not treated as independent people. The six Spearman
447 coefficients have no physiological p-values or confirmatory testing family. They are separate from
448 Q15's frozen two-cohort tests and cannot retrospectively confirm the decoder contrasts.

449 **Reproducibility and Validation Scope**

450 Saved predictions were independently reaggregated to check trial identities, scoring populations,
451 person-level confusion matrices, seed-averaged scores and reported statistics. This numerical
452 reanalysis neither loaded decoder checkpoints nor generated predictions. The Cho2017/Lee2019
453 experiment also underwent a completed replay from raw data to epochs and frozen-model
454 predictions; the earlier PhysioNet portability audit checked saved outputs without replaying
455 inference. Validation of earlier experiments remains experiment-specific, as documented in the
456 supplement.

457 The physiological analysis was checked independently against native BNCI recordings. An
458 explicit detrending, periodic-Hann FFT and trapezoidal integration calculation reproduced all
459 31,104 C3/Cz/C4 trial-band pairs: the maximum relative power difference was 1.85×10^{-15} and
460 the maximum absolute dB difference was 7.11×10^{-15} , within fixed 10^{-10} tolerances. Raw
461 spectral replay covered these three central channels; identities and eligibility were checked for all
462 228,096 saved rows, and all aggregation and association tables were recomputed. Numerical
463 agreement establishes implementation consistency, not independent laboratory replication.
464 OpenAI Codex assisted analysis-code development and verification-code preparation. The
465 computations use explicitly specified algorithms; the archived validation reports document the
466 numerical checks. AI-assisted code generation is distinct from a learned decoder or a physiological
467 model.

468 **Ethics statement**

469 This study involved secondary analysis of public EEG recordings and recruited no new
470 participants. The Cho2017 and Lee2019 source publications report collection ethics review and
471 written informed consent [7,9]. Those collection approvals are not approvals for this secondary
472 analysis. The author confirms that neither ethics approval nor an exemption was required for this
473 secondary analysis.

474 **Data and code availability**

475 Public BNCI2014_001, PhysioNet, Cho2017 and Lee2019 recordings are available from the
476 original providers [1,2,5,6,7,8,9,10], subject to their licenses and terms of use. PhysioNet platform
477 publications are cited separately [27,28]. Analysis code, saved predictions, participant-level
478 summaries, parameter specifications and numerical validation reports are available at
479 <https://github.com/jackzhu119/cross-subject-mi-eeeg>. The supplementary reproducibility
480 documentation identifies the exact source and result revisions. Raw EEG is not redistributed in the
481 manuscript bundle.

482 **AI assistance and verification**

483 ChatGPT (GPT-6, as reported by the author) assisted information retrieval, grammar revision, and
484 preparation of instructions for publishing research materials to GitHub. OpenAI Codex assisted
485 source and literature checks, analysis-code and verification-code development, saved-result
486 consistency checks, figure preparation, and manuscript drafting and language revision. The author
487 personally reviewed and revised the final manuscript and takes responsibility for its scientific
488 content.

489 The underlying Codex model version is not reliably documented in the project records. Saved-
490 result checks compared reported numbers with frozen outputs and archived independent validation
491 reports; these checks did not rerun models or physiology experiments during manuscript
492 preparation. The author reports personal review and revision of the scientific content. No AI tool is
493 an author.

494 **Results**

495 **Traditional baselines and heterogeneous duration selection**

496 The all-trial four-class broad CSP8 + LDA baseline achieved 39.04% balanced accuracy; broad
497 CSP + SVM and the full filter-bank LDA comparator achieved 37.50% and 38.12%. EEGNet with
498 duration selected by raw mean loss achieved 33.72%, below the broad CSP baseline. These
499 descriptive comparisons do not establish superiority of a model class. The mean-loss schedule
500 selected only one or two epochs for S2, S3, and S8, while source-validation optima varied
501 substantially across folds.

502 Source-only per-channel normalization increased mean balanced accuracy to 37.93%, a +4.21 pp
503 difference. The archived paired 95% interval spans -1.12 to $+13.49$ pp, whereas the median
504 change was only $+0.23$ pp. Effects varied across participants, including negative transfer in S8, so
505 the average gain does not establish a general normalization benefit. Selected durations also

506 changed for S3 (2 to 16 epochs) and S5 (13 to 12); the population contrast therefore combines
507 normalization and duration changes.

508 The S3 diagnostic separates normalization from training duration. With two epochs, raw input
509 gave 26.10% balanced accuracy and source-normalized input gave 25.69%. With 16 epochs, raw
510 and normalized inputs gave 67.71% and 65.86%. The short-run raw condition assigned
511 approximately 94% of predictions to one class. Extending training reproduced the recovery for this
512 participant, but this four-cell diagnostic does not establish that duration explains all cross-person
513 failures.

514 Mean-rank selection increased equal-person mean balanced accuracy to 42.67%, a +8.95 pp
515 change over raw mean-loss selection. The archived bootstrap interval is [+1.05,+19.96] pp, and the
516 median change is +1.85 pp. Seven participants improved and two were unchanged. S3 and S8
517 account for 87.57% of the summed improvement. The average gain therefore reflects a
518 concentrated benefit; prior development on the same dataset and disagreement among small-
519 sample statistics limit its generalizability.

520 The corresponding saved summary is presented in Table 4.

521 Table 4. Selected completed four-class conditions; all nine participants.

Condition	Bands (Hz)	Mean BA (%)	Interpretation
Broad CSP8 + LDA	4–40	39.04	Single source fit/fold
Filter-bank CSP + LDA	4–40 bank	38.12	Fixed filter-bank control
EEGNet, mean loss	4–40	33.72	Source-selected duration
EEGNet, source norm	4–40	37.93	Source-only fitted scale
EEGNet, mean rank	4–40	42.67	Reused source curves
Shared mu/beta	8–13 / 13–30	42.19	Representation-specific rank
Broad, fixed 20	4–40	43.99	Later runtime

Condition	Bands (Hz)	Mean BA (%)	Interpretation
Broad, matched CE schedule	4–40	33.92	Same runtime as fixed 20

522 The primary population includes all trials. Neural scores are averaged over three fixed seeds within
523 person before population summaries. Rows span different experimental batches, so runtime
524 variation is controlled only in explicitly matched comparisons. The table summarizes declared
525 controls rather than a target-selected ranking; full representation and robustness inventories are in
526 the supplement.

527 The corresponding saved summary is presented in Fig 2.

528 Fig 2. Training-duration selection and participant-level internal results. A: archived source-
529 selected epochs. B: each participant’s all-trial balanced accuracy averaged over three fixed seeds;
530 the dotted line marks four-class chance (25%). C: the S3 four-cell diagnostic, comparing raw and
531 source-normalized input at two fixed durations. Error bars show standard deviations across three
532 seeds, not uncertainty across people. This diagnostic concerns one person.

533 **Matched-runtime duration and source-composition controls**

534 In the additive matched-runtime experiment, the historical mean-loss duration schedule gave
535 33.92% balanced accuracy, compared with 43.99% for fixed 20 epochs. The fixed-minus-
536 historical-schedule mean was +10.06 pp, with a bootstrap 95% interval of [+0.68,+21.53] pp and
537 exact sign-flip $p = 0.125$. Four of nine people improved and five worsened; the median difference
538 was -0.174 pp. The positive mean thus reflects large recoveries rather than a typical-person
539 benefit. The average schedule difference persisted under matched runtime, but it does not establish
540 20 epochs as universally preferable: historical epoch counts were reused rather than newly selected
541 under that runtime.

542 Mean balanced accuracy increased descriptively with source count. Broad fixed-duration models
543 gave 32.96%, 38.89%, 43.16%, and 43.99% at $k = 2, 4, 6$, and 8; shared models gave 34.74%,
544 38.89%, 42.48%, and 43.98%. Performance varied substantially with source identity at each count.

545 The two-source broad-minus-eight-source difference was -11.03 pp, but its Holm-adjusted
546 exploratory p was 0.164. No source-count/selection/session family showed a corrected significant
547 result. Contrasts between the two sessions were small and heterogeneous. These trajectories are
548 consistent with sensitivity to source training information, but do not separate source diversity from
549 sample count or optimizer-update budget.

550 The corresponding saved summary is presented in Fig 3.

551 Fig 3. Source composition and matched-runtime controls. A: thick lines show equal-person means;
552 thin lines show individual target-person trajectories averaged over the four predefined identity
553 subsets and three seeds at $k = 2, 4, 6$. Every source count uses 20 epochs; $k = 8$ reuses the fixed-20
554 anchor. B: all nine fixed-20-minus-historical-CE-schedule effects in the matched runtime. The
555 interval reverses the sign of the archived matched contrast. Source count also changes data volume
556 and optimizer updates, so these exploratory sensitivities do not isolate source count alone.

557 **Spectral, spatial, and architectural interventions**

558 The internal shared mu/beta condition achieved 42.19% balanced accuracy, compared with
559 42.67% for the broad mean-rank baseline, a -0.48 pp mean difference. The archived Q10-V001
560 exploratory paired interval is $[-3.21, +1.41]$ pp (20,000 draws, seed 20260924), with exact sign-
561 flip $p = 0.914$. This comparison did not demonstrate a spectral-sharing gain. The restricted broad
562 band and individual bands provide descriptive frequency controls, while fixed-duration
563 representation controls retain their declared schedules.

564 The source-only CSP8 and PCA8 neural projections achieved 41.48% and 42.06% balanced
565 accuracy. Architectural controls achieved 42.99% for independent two-band networks, 42.19% for
566 the broadband capacity comparator, 42.66% for early stacking, and 39.80% for four-band sharing.
567 These alternatives provide no evidence of a general shared-band advantage. Differences in
568 preprocessing, projected width, BatchNorm behaviour, and compute also preclude attributing the
569 contrasts to a single physiological factor. Full model means and participant-level results
570 accompany the manuscript.

571 Source-only robustness has no corrected established gain

572 Source-pooled whitening, source-balanced ERM, and gain perturbation had positive point
573 estimates relative to the historical broad rank baseline, whereas dropout conditions were lower. No
574 intervention established an improvement after correction. The matched GroupDRO-minus-
575 balanced-ERM difference was -6.45 pp (unadjusted bootstrap interval $[-11.47, -1.79]$ pp; exact
576 sign-flip $p = 0.046875$, Holm-adjusted $p = 0.140625$). This adverse estimate warrants follow-up
577 but does not establish that distributionally robust optimization generally harms EEG decoding. The
578 whitening and gain perturbation intervals include zero, and comparison with a historical model
579 from a different runtime further limits attribution.

580 The corresponding saved summary is presented in Table 5.

581 Table 5. Exploratory source-only robustness (nine participants).

Condition	BA (%)	Control	Δ (pp)	95% CI (pp)	Holm p
Whitening	44.34	Q8 historical	+1.67	$[-0.39, +3.67]$	0.3594
Balanced ERM	43.18	Q8 historical	+0.51	$[-0.71, +1.77]$	0.4766
GroupDRO	36.73	Balanced ERM	-6.45	$[-11.47, -1.79]$	0.1406
Channel dropout	41.00	Q8 historical	-1.67	$[-4.51, +0.64]$	0.7617
Gain perturbation	43.42	Q8 historical	+0.75	$[-0.34, +1.86]$	0.7617
Dropout + gain	40.92	Q8 historical	-1.75	$[-5.05, +0.91]$	0.7617

582 Intervals are the archived, unadjusted person-bootstrap intervals; p-values use the declared within-
583 family Holm correction. Every comparison with Q8 crosses runtimes. GroupDRO versus balanced
584 ERM matches batching within the same runtime. Neither intervals nor p-values are used to select a
585 model on target outcomes.

586 Earlier PhysioNet transfer: a small, uncertain primary effect

587 In binary BNCI development, broad, shared, and CSP models achieved 68.21%, 67.35%, and
588 61.54% balanced accuracy. These development outcomes cannot be pooled with the four-class

analyses. The separately frozen all-nine-source models then generated predictions for all 109 external subjects. Final validation confirmed 4,918 trial identities and 34,426 saved prediction rows. No model was fitted to target data.

External equal-person balanced accuracy was 61.81% for broad EEGNet, 62.39% for shared μ/β , and 54.46% for CSP4 + LDA. The predeclared shared-minus-broad primary difference was +0.573 pp, with a 95% bootstrap interval of [-0.097, +1.242] pp. Sixty-three people favoured shared, 43 favoured broad, and three tied; the exact tie-excluding sign-test p was 0.06446. The interval includes zero, and the sign test does not establish external superiority.

Individual performance varied widely. The shared model ranged from 39.33% to 97.07% balanced accuracy, with a median of 58.73%. Differences averaged across three seeds also varied in direction across people. The three resampled subjects had mean shared-minus-broad $\Delta = +0.514$ pp, compared with +0.575 pp for the 106 native-160-Hz subjects. The three-person subgroup comparison is descriptive and has no powered subgroup test. The result incorporates the sampling-rate and validation-portability amendments and therefore does not represent an unmodified execution of the original external protocol.

The corresponding saved summary is presented in Fig 4.

Fig 4. Separate-cohort binary transfer. A: empirical cumulative distributions of three-seed person-mean balanced accuracy, with one CSP score per person; the dotted vertical line marks binary chance (50%). B: all 109 paired primary effects, with zero and the mean indicated. The interval and p-value are the frozen person-bootstrap and tie-excluding sign-test outputs. Three participants were deterministically resampled under the disclosed metadata amendment.

The corresponding saved summary is presented in Table 6.

Table 6. Separate binary external evaluation (109 participants).

Frozen model	Mean BA (%)	Subject SD (pp)	Median BA (%)
Broad EEGNet	61.81	12.56	58.23

Frozen model	Mean BA (%)	Subject SD (pp)	Median BA (%)
Shared mu/beta	62.39	12.92	58.73
CSP4 + LDA	54.46	7.90	52.08

Deep-model scores are averaged over seeds within person. The primary shared-minus-broad contrast is +0.573 pp, 95% CI [-0.097,+1.242], $p = 0.06446$. CSP provides context and is outside the primary hypothesis. The Methods describe the sampling-rate amendment and versioned portability validation.

Completed Cho2017 and Lee2019 transfer: lower balanced accuracy on Cho2017 and an uncertain Lee2019 contrast

The completed Q15 outputs cover 21,320 unique external trials and 149,240 saved model/seed prediction rows, seven per trial. Independent checks verified canonical labels and event identities, person-level confusion matrices, aggregate balanced accuracies, frozen statistical outputs, and links to the declared artifacts. Original source fits remained 15, with zero new migration fits and zero target fits. These counts confirm source-only fitting; they do not constitute independent repetitions of the external experiment.

On Cho2017, equal-person balanced accuracy was 59.82% for broad EEGNet and 58.31% for shared mu/beta EEGNet. The primary shared-minus-broad mean was -1.513 pp (95% person-bootstrap interval [-2.249,-0.804] pp), with two-sided Monte Carlo sign-flip $p = 0.000049998$ and two-cohort Holm-adjusted $p = 0.000099995$. The raw p-value is the Monte Carlo procedure's minimum attainable value and does not resolve probabilities below that value. Twelve people favoured shared, 37 favoured broad, and three were tied within a 10^{-15} numerical tolerance. Shared performed worse in this declared complete-pipeline comparison after correction.

On Lee2019, broad and shared models achieved 65.51% and 65.72% balanced accuracy. The primary difference was +0.204 pp (95% interval [-0.515,+0.969] pp); raw and Holm-adjusted p-values were both 0.60367. Twenty-three people favoured shared, 30 favoured broad, and one tied. This small, uncertain contrast establishes neither a reliable improvement nor equivalence. The two

635 sessions were combined within each person, so neither the 108 session records nor repeated model
636 seeds increase the primary sample size beyond 54.

637 The corresponding saved summary is presented in Table 7.

638 Table 7. New external binary evaluation under the Q15 freeze.

Cohort / model	BA (%)	Person SD (pp)	Left recall (%)	Right recall (%)
Cho2017 / Broad EEGNet	59.82	8.64	46.79	72.86
Cho2017 / Shared mu/beta	58.31	7.22	37.52	79.10
Cho2017 / CSP4 + LDA	51.99	5.67	94.10	9.88
Lee2019 / Broad EEGNet	65.51	10.97	75.56	55.46
Lee2019 / Shared mu/beta	65.72	10.76	72.69	58.75
Lee2019 / CSP4 + LDA	52.77	5.57	85.98	19.56

639 Neural scores and class recalls are averaged over three seeds within person, then equally over
640 people. Person SD describes heterogeneity, not a standard error. Lee includes only offline-training
641 runs, pooled across two sessions. The primary shared-minus-broad contrasts are -1.513 pp
642 $[-2.249, -0.804]$, Holm $p = 0.000100$ for Cho, and $+0.204$ pp $[-0.515, +0.969]$, Holm $p = 0.604$ for
643 Lee. Class recalls and CSP comparisons are descriptive; no additional confirmatory tests were
644 added.

645 The CSP comparator was near binary chance on both cohorts and strongly favoured left-hand
646 labels. On Cho2017, mean left-hand recall fell from 46.79% with broad input to 37.52% with
647 shared input, while right-hand recall increased from 72.86% to 79.10%. On Lee2019, the
648 corresponding recalls changed from 75.56% to 72.69% and from 55.46% to 58.75%. Balanced

649 accuracy summarizes these asymmetric shifts but does not show which class gains or loses recall.
650 These descriptive checks used the frozen outputs and did not trigger any threshold, label, adapter,
651 or model change.

652 The corresponding saved summary is presented in Fig 5.

653 Fig 5. Frozen external evaluation on Cho2017 and Lee2019. Frozen Q15 external evaluation.

654 Cohorts are shown separately: person/model scores, equal-person means and person standard
655 deviations, and frozen paired person-bootstrap 95% intervals. Deep-model scores average three
656 seeds within person; Lee sessions are pooled before person scoring. Holm p-values refer to the
657 predeclared two-cohort shared-minus-broad family. The common 21-channel, two-second pipeline
658 differs from Q14, and no pooled cross-cohort test was specified.

659 **Descriptive BNCI Sensor Physiology and Decoder Associations (Q16)**

660 All 18 raw files, 108 labelled runs and 5,184 trials passed the fixed coverage checks. All 228,096
661 trial-channel-band rows had finite, positive baseline and task power. No trial, session or person
662 was excluded; the 488 provider-flagged trials remained in the primary summaries. Fig 6 shows
663 equal-person scalp means and all nine C3/Cz/C4 profiles. Task/baseline changes calculated with
664 the fixed estimator varied by sensor, imagined hand and person.

665 Fig 6. Descriptive native BNCI task/baseline power changes. A–D: equal-person hand-imagery
666 means for mu (8–13 Hz) and beta (13–30 Hz) at all 22 EEG sensors. In the schematic
667 standard-1020 geometry, anterior is upward and left scalp is on the viewer’s left; linear
668 interpolation is restricted to the sensor convex hull. No sensor or interpolated grid value is clipped
669 by the fixed ± 6 -dB range. E–F: all nine equal-session C3/Cz/C4 profiles (thin lines) and equal-
670 person means (thick lines). Trial dB uses cue baseline $[-1.5, -0.5]$ s and task $[0.5, 2.5]$ s, with one
671 versus three overlapping Welch segments. Negative values indicate lower task power under this
672 estimator; unequal window precision can bias log ratios. All trials and artifact flags are retained.
673 Template topography neither localizes cortical sources nor attributes decoder decisions.

674 Mu signed laterality averaged -0.250 dB, with six negative and three positive person descriptors;
675 beta averaged -0.168 dB, with eight negative and one positive descriptor (Table 8). Contralateral
676 and ipsilateral group means were both negative in each band, but absolute contralateral changes
677 were not negative for every person. For example, S2 had negative mu laterality (-0.203 dB) even
678 though both its contralateral ($+1.491$ dB) and ipsilateral ($+1.694$ dB) means were positive.
679 Separate session, hand and flagged/unflagged summaries are available in the verified
680 supplementary tables as descriptive sensitivity checks; they do not replace the all-trial population.
681 Table 8. Baseline-relative hand power and signed C3/C4 laterality.

Band	Contra. mean (dB)	Ipsi. mean (dB)	Mean L (dB)	Person L range (dB)	L signs (-/+)
Mu, 8–13 Hz	-0.308	-0.058	-0.250	-1.432 to $+0.336$	6 / 3
Beta, 13–30 Hz	-0.439	-0.272	-0.168	-0.519 to $+0.133$	8 / 1

682 Each of nine people receives equal weight, using equal-session mean trial dB from both hands.
683 Contra. and ipsi. average the corresponding C3/C4 hand components; L is contra. minus ipsi. All
684 2,592 hand trials, including 246 flagged trials, are retained. Negative L means a lower contralateral
685 than ipsilateral task/baseline change; it does not establish an absolute contralateral power decrease.
686 Values are rounded to three decimals here and retained at full precision in the data tables. These
687 fixed-estimator ratios are descriptive, not unbiased ERD estimates.
688 For the designated broad-EEGNet endpoint, descriptive associations between signed laterality and
689 saved binary LOSO balanced accuracy were $\rho = -0.150$ for mu and $+0.600$ for beta. Shared-
690 input coefficients were -0.083 and $+0.483$, and CSP4+LDA coefficients were $+0.067$ and $+0.317$,
691 respectively (Table S4; S1 Appendix, Fig S3). All six coefficients use the same nine people and
692 have no p-values. The positive beta coefficient indicates that higher balanced accuracy
693 accompanied a larger, less-negative signed descriptor; it does not support a benefit from stronger
694 relative contralateral suppression.

695 Discussion

696 Training-duration Sensitivity in a Small Development Benchmark

697 Training duration was a major source of internal performance variation. At two epochs, S3
698 remained near chance with either raw or source-normalized inputs; at 16 epochs, both performed
699 substantially better. Rank aggregation changed the source-selected schedule and recovered
700 performance for S3 and S8. The matched-runtime control also retained a large mean difference
701 between fixed training and the historical mean-loss schedule. An apparent cross-person model
702 collapse should therefore be examined alongside selected duration, source-validation trajectories,
703 per-class recall and prediction concentration. These outputs help distinguish an inadequate
704 stopping schedule from a failure attributed prematurely to the input representation.

705 These findings do not establish rank selection as a general solution. Within-fold ranks remove
706 loss-scale differences but discard the size of those differences; equal fold weighting imposes a
707 further aggregation choice. The candidate was developed after this benchmark had been inspected.
708 In addition, the matched-runtime fixed-duration condition improved the mean while the median
709 paired difference was negative. The external evaluations did not compare rank with mean-loss
710 selection and cannot independently confirm the internal gain. A prospective test would fix the
711 competing selection rules before observing new-person outcomes and match runtime, source
712 examples and update budgets.

713 Complete Spectral Pipelines Have Cohort-dependent Outcomes

714 The frozen evaluations limit claims of a consistent shared-input advantage. Shared μ/β input
715 gave no internal advantage, a small uncertain positive estimate on PhysioNet, a corrected adverse
716 contrast on Cho2017 and an uncertain near-zero contrast on Lee2019. The Cho2017 finding
717 contradicts a consistent advantage of this fixed shared-input pipeline under the evaluated source-
718 only adapter. It does not show that μ/β rhythms lack information, that parameter sharing is
719 generally harmful or that a different band definition would perform better.

720 The class recalls reveal a trade-off hidden by the cohort means. On Cho2017, the shared-input
721 pipeline reduced mean left-hand recall while increasing right-hand recall; on Lee2019, opposing
722 recall changes partly offset one another. These patterns describe the frozen predictions. They do
723 not justify correcting labels after inspection or attributing the contrast to physiology. Montage,
724 acquisition reference, amplitude scale, temporal alignment and band-logit averaging are possible
725 contributors, but the study did not isolate any of them.

726 One shared parameter set does not imply a cheaper computation: the network processes two band
727 views and pools their BatchNorm exposure. Independent-branch controls alter both capacity and
728 normalization, and the broadband capacity comparator retains a small disclosed parameter
729 mismatch. The source-fitted CSP model also transfers poorly and asymmetrically in Q15.

730 Performance above this contextual comparator does not establish practical or clinical reliability.
731 Contemporary adaptation and foundation-model methods were not compared under a common
732 budget, so the results support no state-of-the-art claim.

733 The external schedules are part of the pipeline contrast: broad/shared fits used 18/17 epochs in
734 Q14 and 14/19 in Q15. Paired differences therefore combine the selected training durations, input
735 views and shared normalization. They do not identify the causal contribution of parameter sharing
736 alone.

737 **Person-level Evidence and Descriptive Signal Characterization**

738 Additional seeds assess sensitivity to stochastic fitting; they do not increase the number of
739 participants. Lee2019 sessions and source-identity subsets likewise provide repeated observations
740 within a person. The paired person-level analysis retains all declared participants and reports the
741 direction of each difference. Its bootstrap intervals condition on the frozen source models and omit
742 uncertainty from another nine-person source cohort, acquisition adapter or development history.

743 For the internal study, overlapping outer training sets further limit an independent-person
744 interpretation.

745 A few large recoveries can improve a mean without producing consistent benefit across people, so
746 tests of average change and directional consistency need not agree. The original rank-selection
747 sign calculation retained exact ties; excluding them conventionally gives a different result from
748 both the archived calculation and the paired t-test. The discrepancy must remain visible rather than
749 be hidden under an incorrect test label. The later internal families retain their corrected null
750 findings. Percentile bootstrap intervals and randomization tests also need not agree in a small
751 sample; an interval excluding zero does not override a corrected null test. Q15 uses a separately
752 frozen two-cohort family: Cho2017 survives correction, whereas Lee2019 establishes neither
753 benefit nor equivalence. These separate tests do not establish a cohort-by-pipeline interaction.
754 PhysioNet has a different archived primary test and an amended protocol; its p-value belongs to
755 that evaluation rather than the newer family.

756 The BNCI signals describe measured sensor activity, not the cause of the decoder contrasts. Mean
757 signed hand laterality was negative in both bands: six of nine people had negative mu laterality
758 and eight had negative beta laterality, with positive signs in the others. A negative relative
759 descriptor can still accompany increased power on both sides, as in S2, so its absolute components
760 must also be examined. The positive broad-model beta correlation means higher balanced accuracy
761 accompanied less-negative relative laterality, rather than stronger suppression. None of the six
762 associations supports a subgroup threshold or a model-selection rule.

763 Neither a lateralized scalp mean nor its correlation with LOSO balanced accuracy locates a cortical
764 source, excludes ocular or muscle contributions, or demonstrates that EEGNet used that signal.
765 This BNCI-only physiology analysis was specified after decoder outcomes were known. It cannot
766 explain the adverse Cho contrast or validate an external physiological mechanism.

767 **Limitations**

768 Successive development conditions used the same nine BNCI people. Excluding the target from
769 each fit does not remove outcome-informed candidate design or dependence through overlapping
770 LOSO training sets. The internal rank-selection gain was concentrated in two people, and the

external evaluations did not test that selection-rule contrast. Controls also differ in runtime, training examples, update exposure, sensor width or approximate capacity. There is no common-budget comparison with contemporary adaptation or foundation-model methods.

Q14 and Q15 differ in montages, windows, references, filtering scope and source checkpoints. Q14's metadata amendment followed partial execution. In the frozen transfer experiment, repository freezes preceded fitting or inference but followed metadata inspection; they are not independently registered preregistration. Cross-provider participant identities remain unverified. The external bootstrap intervals condition on one frozen source cohort and omit source-sampling, adapter and development-history uncertainty.

The external duration schedules were selected using four fixed BNCI validation groups with unequal participant counts (2/2/2/3) and equal fold weighting. The grouping was frozen before target inference, but it is one arbitrary source-validation design. Its sensitivity was not tested prospectively.

Q15 used a common operational tensor and a committed adapter for transfer across three providers. Matching tensor dimensions does not make the physical measurements equivalent. Native export voltage calibration, the original Cho acquisition reference and hardware cue latency were not independently verified. The Cho channel order was mapped from the publication montage because the MATLAB payload has no channel-name strings. Lee event indexing retains a one-native-sample convention discrepancy. Numerical consistency checks cannot resolve these acquisition uncertainties.

Q14 and Q15 also differ in the extent of verification. The earlier PhysioNet evaluation required a sampling-rate amendment after partial execution; its later portability check examined saved outputs without replaying inference. Q15 froze preprocessing before source fitting and the inference contract before external prediction, and its completed scientific validator reports independent raw-to-prediction replay. The manuscript checks recalculate saved outputs and verify

796 hash bindings. These provide different levels of reproducibility and should not be treated as
797 interchangeable.

798 Evaluation was offline and used zero-phase filtering. Online latency, prospective assistive control,
799 clinical populations and patient benefit were not tested. The two-second Q15 crop omits part of
800 each imagery interval and differs from the providers' original benchmarks. A subsequent study
801 could preregister a calibrated adapter and frozen source-only models for another cohort, specifying
802 per-class and per-person outcomes before access. Retuning on the present held-out cohorts would
803 be further development, not independent external confirmation.

804 The physiology estimates compare a one-segment baseline with a three-segment task estimator;
805 unequal log-power variance can bias the dB difference. Negative signed laterality can arise from
806 ipsilateral power increases without absolute contralateral suppression. The descriptor and decoder
807 use different task windows, and nine descriptive person records with overlapping LOSO fits
808 cannot establish a subgroup biomarker or mechanistic attribution. Baseline-relative ratios cancel
809 only a constant multiplicative gain; they do not correct reference, timing, spatial mapping or
810 artifacts. External voltage calibration remains unknown, and external physiology has not been
811 evaluated.

812 **Conclusions**

813 Training-duration selection was consequential in this nine-person BNCI development benchmark,
814 but its gains were concentrated rather than broadly shared. Fixed shared-input pipelines showed no
815 consistent external advantage: balanced accuracy was 1.513 pp lower than broad EEGNet on
816 Cho2017, and the Lee2019 estimate remained compatible with small effects in either direction.
817 Unequal source-selected durations are part of these complete-pipeline comparisons, which do not
818 isolate parameter sharing. The post-outcome BNCI sensor analysis found negative mean signed
819 hand laterality in mu and beta with substantial participant variation; it established neither a

820 decoder mechanism nor external physiological replication. Unknown acquisition calibration and
821 offline evaluation limit deployment conclusions.

822 **Acknowledgments**

823 No additional acknowledgments are reported.

824 **References**

- 825 1. Brunner C, Leeb R, Müller-Putz GR, Schlögl A, Pfurtscheller G. BCI Competition 2008 – Graz
826 data set A. Official competition dataset description. 2008. Available from:
827 https://www.bbci.de/competition/iv/desc_2a.pdf [cited 2026 Oct 8].
- 828 2. Tangermann M, Müller KR, Aertsen A, Birbaumer N, Braun C, Brunner C, et al. Review of the
829 BCI Competition IV. *Front Neurosci.* 2012;6. doi:10.3389/fnins.2012.00055
- 830 3. Jayaram V, Barachant A. MOABB: trustworthy algorithm benchmarking for BCIs. *J Neural*
831 *Eng.* 2018;15(6):066011. doi:10.1088/1741-2552/aadea0
- 832 4. Lotte F, Bougrain L, Cichocki A, Clerc M, Congedo M, Rakotomamonjy A, et al. A review of
833 classification algorithms for EEG-based brain–computer interfaces: a 10 year update. *J Neural*
834 *Eng.* 2018;15(3):031005. doi:10.1088/1741-2552/aab2f2
- 835 5. Schalk G, McFarland DJ, Hinterberger T, Birbaumer N, Wolpaw JR. EEG Motor
836 Movement/Imagery Dataset. dataset. 2009. doi:10.13026/C28G6P
- 837 6. Schalk G, McFarland DJ, Hinterberger T, Birbaumer N, Wolpaw JR. BCI2000: A General-
838 Purpose Brain-Computer Interface (BCI) System. *IEEE Trans Biomed Eng.* 2004;51(6):1034-
839 1043. doi:10.1109/TBME.2004.827072
- 840 7. Cho H, Ahn M, Ahn S, Kwon M, Jun SC. EEG datasets for motor imagery brain–computer
841 interface. *GigaScience.* 2017;6(7). doi:10.1093/gigascience/gix034
- 842 8. Cho H, Ahn M, Ahn S, Kwon M, Jun SC. Supporting data for "EEG datasets for motor imagery
843 brain computer interface". dataset. 2017. doi:10.5524/100295

844 9. Lee MH, Kwon OY, Kim YJ, Kim HK, Lee YE, Williamson J, et al. EEG dataset and OpenBMI
845 toolbox for three BCI paradigms: an investigation into BCI illiteracy. *GigaScience*. 2019;8(5).
846 doi:10.1093/gigascience/giz002

847 10. Lee MH, Kwon OY, Kim YJ, Kim HK, Lee YE, Williamson J, et al. Supporting data for "EEG
848 Dataset and OpenBMI Toolbox for Three BCI Paradigms: An Investigation into BCI Illiteracy".
849 dataset. 2019. doi:10.5524/100542

850 11. Gulrajani I, Lopez-Paz D. In Search of Lost Domain Generalization. *International Conference*
851 *on Learning Representations*. 2021. Available from: <https://iclr.cc/virtual/2021/poster/2998> [cited
852 2026 Oct 8].

853 12. Ramoser H, Muller-Gerking J, Pfurtscheller G. Optimal spatial filtering of single trial EEG
854 during imagined hand movement. *IEEE Trans Rehabil Eng*. 2000;8(4):441-446.
855 doi:10.1109/86.895946

856 13. Ang KK, Chin ZY, Wang C, Guan C, Zhang H. Filter Bank Common Spatial Pattern
857 Algorithm on BCI Competition IV Datasets 2a and 2b. *Front Neurosci*. 2012;6.
858 doi:10.3389/fnins.2012.00039

859 14. Ang KK, Chin ZY, Zhang H, Guan C. Filter Bank Common Spatial Pattern (FBCSP) in Brain-
860 Computer Interface. 2008 IEEE International Joint Conference on Neural Networks (IEEE World
861 Congress on Computational Intelligence). 2008:2390-2397. doi:10.1109/IJCNN.2008.4634130

862 15. Lawhern VJ, Solon AJ, Waytowich NR, Gordon SM, Hung CP, Lance BJ. EEGNet: a compact
863 convolutional neural network for EEG-based brain-computer interfaces. *J Neural Eng*.
864 2018;15(5):056013. doi:10.1088/1741-2552/aace8c

865 16. Schirrmester RT, Springenberg JT, Fiederer LDJ, Glasstetter M, Eggensperger K,
866 Tangermann M, et al. Deep learning with convolutional neural networks for EEG decoding and
867 visualization. *Hum Brain Mapp*. 2017;38(11):5391-5420. doi:10.1002/hbm.23730

- 868 17. Zhong XC, Wang Q, Liu D, Chen Z, Liao JX, Sun J, et al. EEG-DG: A Multi-Source Domain
869 Generalization Framework for Motor Imagery EEG Classification. *IEEE J Biomed Health Inform.*
870 2025;29(4):2484-2495. doi:10.1109/jbhi.2024.3431230
- 871 18. Zheng Y, Wu S, Chen J, Yao Q, Zheng S. Cross-Subject Motor Imagery
872 Electroencephalogram Decoding with Domain Generalization. *Bioengineering (Basel).*
873 2025;12(5):495. doi:10.3390/bioengineering12050495
- 874 19. Jiang WB, Zhao LM, Lu BL. Large Brain Model for Learning Generic Representations with
875 Tremendous EEG Data in BCI. *The Twelfth International Conference on Learning*
876 *Representations*. 2024. Available from: <https://openreview.net/forum?id=QzTpTRVtrP> [cited 2026
877 Oct 8].
- 878 20. Wang G, Liu W, He Y, Xu C, Ma L, Li H. EEGPT: Pretrained Transformer for Universal and
879 Reliable Representation of EEG Signals. *Advances in Neural Information Processing Systems*.
880 2024;37. doi:10.52202/079017-1239
- 881 21. Pfurtscheller G, Lopes da Silva FH. Event-related EEG/MEG synchronization and
882 desynchronization: basic principles. *Clin Neurophysiol.* 1999;110(11):1842-1857.
883 doi:10.1016/s1388-2457(99)00141-8
- 884 22. Pfurtscheller G, Neuper C. Motor imagery activates primary sensorimotor area in humans.
885 *Neurosci Lett.* 1997;239(2-3):65-68. doi:10.1016/s0304-3940(97)00889-6
- 886 23. Sagawa S, Koh PW, Hashimoto TB, Liang P. Distributionally Robust Neural Networks for
887 Group Shifts: On the Importance of Regularization for Worst-Case Generalization. *International*
888 *Conference on Learning Representations*. 2020. Available from: [https://openreview.net/forum?](https://openreview.net/forum?id=ryxGuJrFvS)
889 [id=ryxGuJrFvS](https://openreview.net/forum?id=ryxGuJrFvS) [cited 2026 Oct 8].
- 890 24. Efron B. Bootstrap Methods: Another Look at the Jackknife. *Ann Stat.* 1979;7(1).
891 doi:10.1214/aos/1176344552
- 892 25. Winkler AM, Ridgway GR, Webster MA, Smith SM, Nichols TE. Permutation inference for
893 the general linear model. *Neuroimage.* 2014;92:381-397. doi:10.1016/j.neuroimage.2014.01.060

894 26. Holm S. A Simple Sequentially Rejective Multiple Test Procedure. Scand J Stat. 1979;6(2):65
895 –70. Available from: <https://www.jstor.org/stable/4615733> [cited 2026 Oct 8].

896 27. Goldberger AL, Amaral LAN, Glass L, Hausdorff JM, Ivanov PC, Mark RG, et al.
897 PhysioBank, PhysioToolkit, and PhysioNet: Components of a New Research Resource for
898 Complex Physiologic Signals. Circulation. 2000;101(23). doi:10.1161/01.CIR.101.23.e215

899 28. Pollard T, Moody BE, Lehman LwH, Gow BJ, Fernandes C, Xie C, et al. PhysioNet as a
900 global platform for biomedical research. Nat Health. 2026;1(8):792-795. doi:10.1038/s44360-026-
901 00096-z

902 **Supporting information captions**

903 S1 Appendix. Supplementary results, descriptive associations, inventories and detailed
904 reproducible methods. Internal supplementary figures and tables retain their original S numbering.

905 S1 Data. Frozen participant-level summaries, saved external predictions, descriptive physiological
906 outputs, source-data provenance and validation evidence. Files are CSV/JSON or gzip-compressed
907 CSV; raw EEG and model weights are excluded.